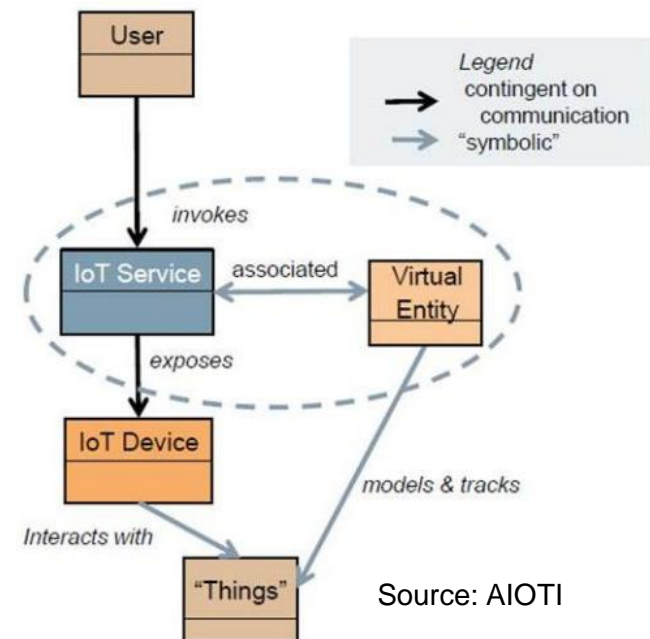


EIOT

MQTT

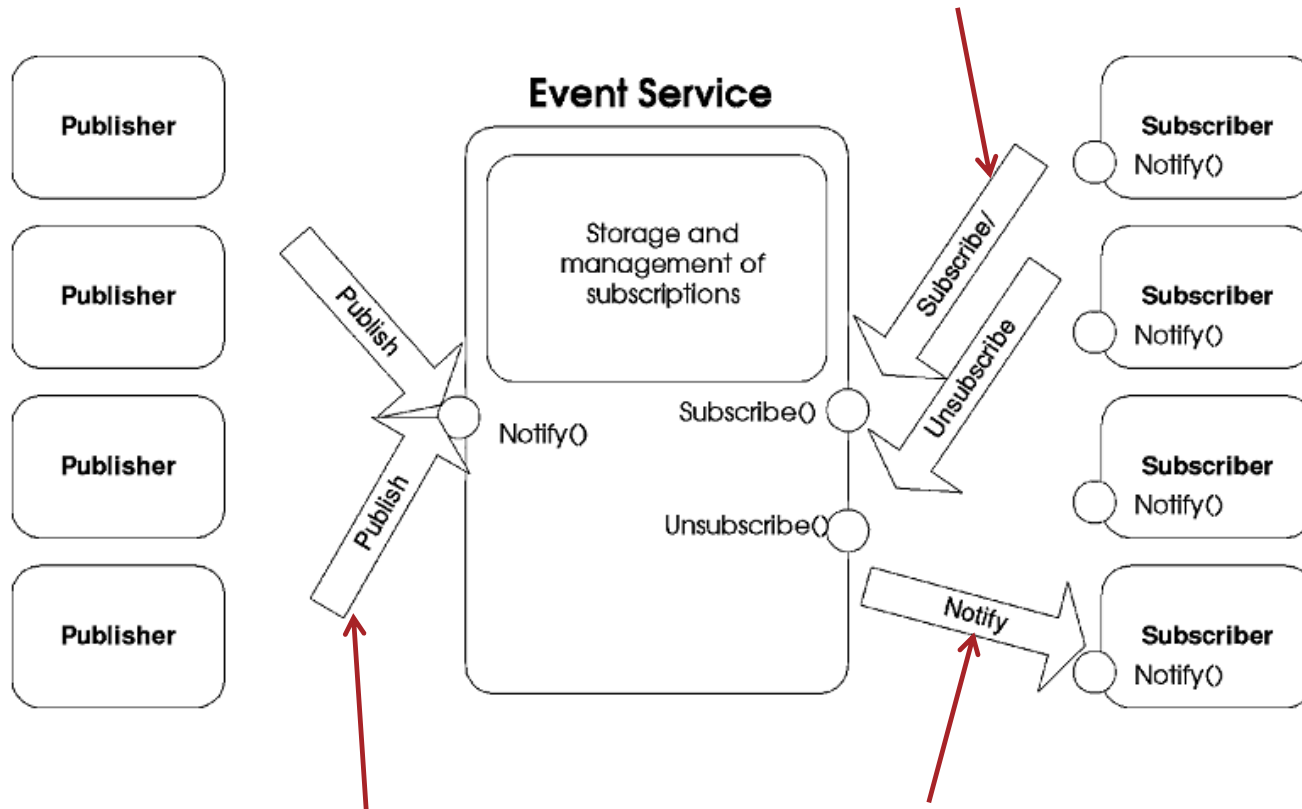
Jarosław Domaszewicz

Instytut Telekomunikacji Politechniki Warszawskiej



PUBLISH/SUBSCRIBE FUNDAMENTALS (1/5)

1. subscribers express interest in selected events



Source:
The Many Faces of Publish/Subscribe
P. Eugster et al.,
ACM Computing Surveys,
Vol. 35, No. 2, 2003



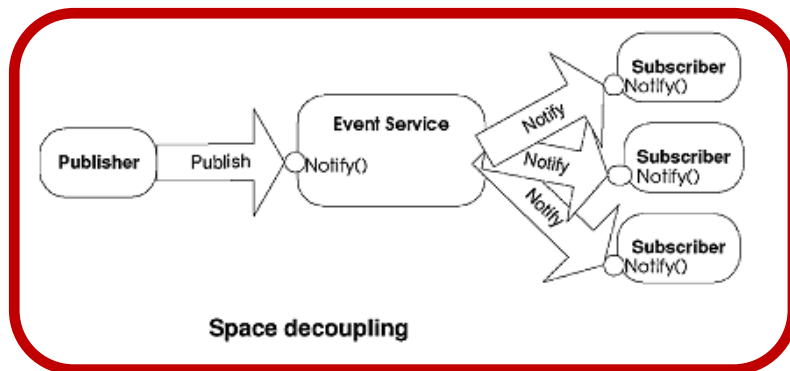
2. publishers produce assorted events 3. subscribers receive events they are interested in

- **publishers, subscribers, event service** (middleware)
- note: push vs. pull, one-to-many, many-to-one

PUBLISH/SUBSCRIBE FUNDAMENTALS (2/5)

- **space decoupling**

- publishers and subscribers do not know about one another (compare with client-server)
- they do need to know about the event service
- publishers do not know how many subscribers participate and vice versa




decoupling is good!

Source: *The Many Faces of Publish/Subscribe*
P. Eugster et al.
ACM Computing Surveys, Vol. 35, No. 2, 2003

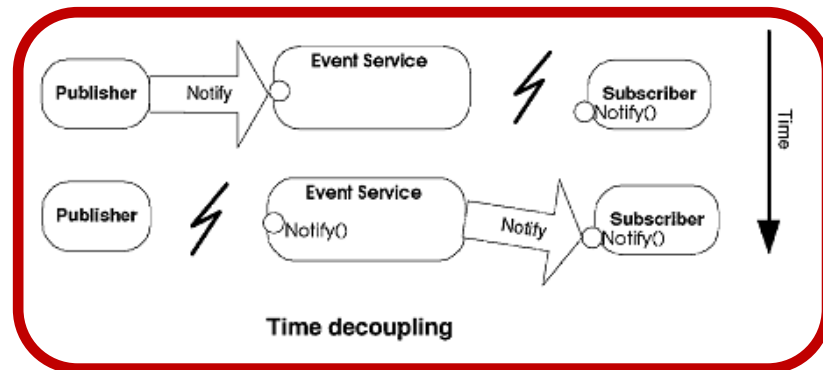
PUBLISH/SUBSCRIBE FUNDAMENTALS (3/5)

- **time decoupling**


- publishing and delivery may occur at different times (compare with client-server) ...
- ... (but most often they occur close in time)



Space decoupling



Time decoupling



Synchronization decoupling


decoupling is good!

Source: *The Many Faces of Publish/Subscribe*
P. Eugster et al.
ACM Computing Surveys, Vol. 35, No. 2, 2003

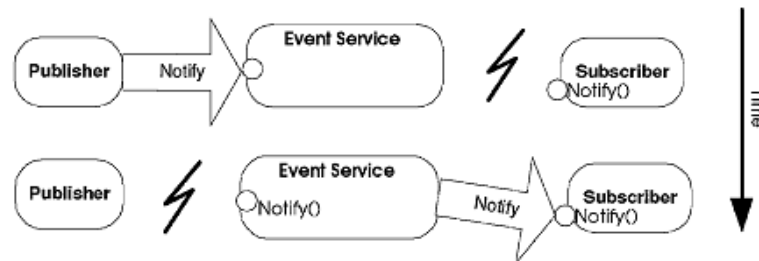
PUBLISH/SUBSCRIBE FUNDAMENTALS (4/5)

- **synchronization decoupling**

- when publishing, publishers are not blocked
- subscribers receive notification asynchronously (via a callback)




Space decoupling



Time decoupling

publishing is non-blocking



Synchronization decoupling

decoupling is good!

Source: *The Many Faces of Publish/Subscribe*
P. Eugster et al.
ACM Computing Surveys, Vol. 35, No. 2, 2003

PUBLISH/SUBSCRIBE FUNDAMENTALS (5/5)

- how to specify events of interest?
- **topic-based** publish/subscribe
 - also called subject-based
 - subscribers subscribe to events published under a given topic
 - publishers tag events with a topic
 - each topic amounts to an individual communication channel
- **content-based** publish/subscribe
 - subscriptions based on the actual content of events
 - a subscriber provides a subscription pattern that refers to the content
- **type-based** publish/subscribe
 - subscription based on the type of (kind of) events

MQTT: KEY FACTS

- an OASIS and ISO standard
 - new: version 5 (but we cover v. 3.1.1)



OASIS Standard

07 March 2019



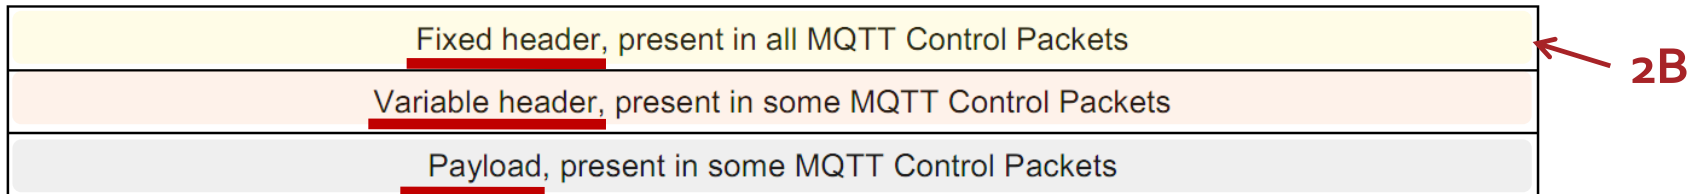
- runs on top of TCP
- implements publish/subscribe
- binary and lightweight
- agnostic as to application data
- some terminology:
 - **MQTT clients:** publishers, subscribers
 - **MQTT broker** (server): event service
 - **MQTT session:** an MQTT-level connection between a client and the broker

MQTT MECHANISMS

- topics and topic filters
- retained messages
- Keep Alive
- will
- persistent sessions
- (application-level) QoS (i.e., reliability)
 - even though it runs on top of TCP
 - takes care of broken TCP connections

MQTT CONTROL PACKET FORMAT

Figure 2.1 – Structure of an MQTT Control Packet



Source:
MQTT Version 3.1.1
OASIS Standard , October 2014

ASIDE: MQTT DATA REPRESENTATIONS

- Bits
 - bits in a byte are labeled 7 through 0; bit number 7 is the most significant bit
- Integer data values (fixed length)
 - 16 bits in big-endian order: the high order byte precedes the lower order byte
- Variable-length integers

- least significant byte first →

1	0 (0x00)	127 (0x7F)
2	128 (0x80, 0x01)	16 383 (0xFF, 0x7F)
3	16 384 (0x80, 0x80, 0x01)	2 097 151 (0xFF, 0xFF, 0x7F)
4	2 097 152 (0x80, 0x80, 0x80, 0x01)	268 435 455 (0xFF, 0xFF, 0xFF, 0x7F)

- UTF-8 strings.

Figure 1.1 Structure of UTF-8 encoded strings

Bit	7	6	5	4	3	2	1	0
byte 1	String length MSB							
byte 2	String length LSB							
byte 3	UTF-8 Encoded Character Data, if length > 0.							

Source: *MQTT Version 3.1.1*
OASIS Standard , October 2014

MQTT CONTROL PACKET FIXED HEADER

Figure 2.2 - Fixed header format

Bit	7	6	5	4	3	2	1	0
byte 1	MQTT Control Packet type				Flags specific to each MQTT Control Packet type			
byte 2...	Remaining Length							

length of variable header and payload

connection establishment

publishing

ensuring QoS
in delivering published
messages

subscription management

heartbeat

disconnection

Name	Value
Reserved	0
CONNECT	1
CONNACK	2
PUBLISH	3
PUBACK	4
PUBREC	5
PUBREL	6
PUBCOMP	7
SUBSCRIBE	8
SUBACK	9
UNSUBSCRIBE	10
UNSUBACK	11
PINGREQ	12
PINGRESP	13
DISCONNECT	14
Reserved	15

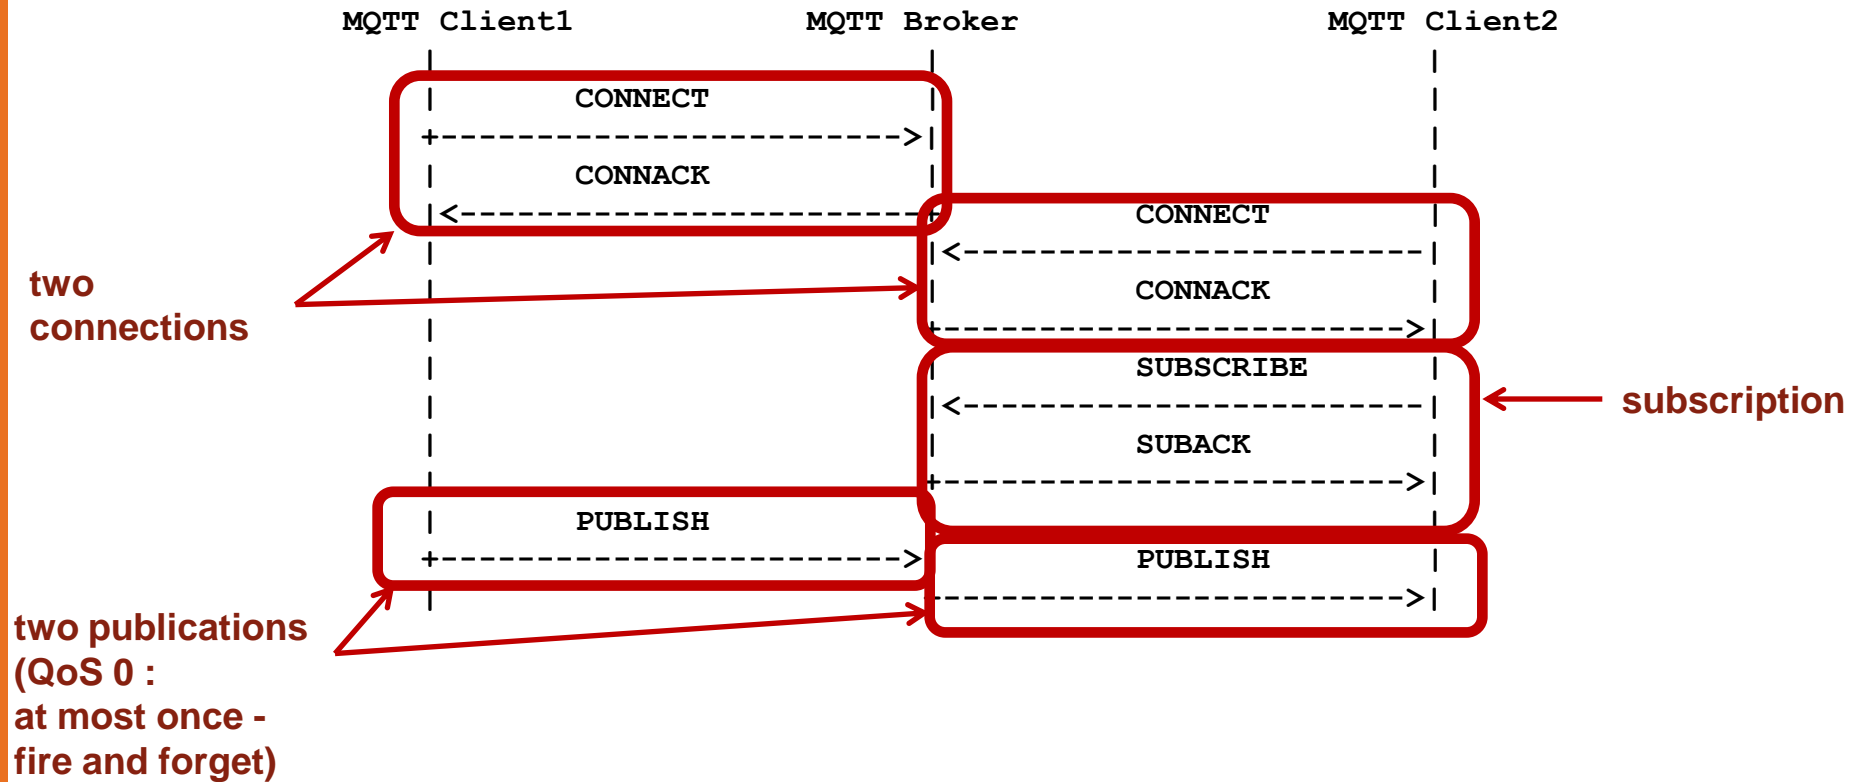
DUP ¹	QoS ²	QoS ²	RETAIN ³
------------------	------------------	------------------	---------------------

flags used in PUBLISH only

Source:
MQTT Version 3.1.1
OASIS Standard , October 2014

client to broker
broker to client

TYPICAL MESSAGE EXCHANGE



MQTT CONTROL PACKET VARIABLE HEADER (1/3)

- In CONNECT:
 - Protocol Name ("MQTT")
 - Protocol Level (rev. level)
 - Connect Flags

Source: *MQTT Version 3.1.1*
OASIS Standard , October 2014

Figure 3.4 - Connect Flag bits

Bit	7	6	5	4	3	2	1	0
	User Name Flag	Password Flag	Will Retain	Will QoS		Will Flag	Clean Session	Reserved
byte 8	X	X	X	X	X	X	X	0

- Keep Alive time
- In CONNACK:
 - CONNECT return code (accepted or refused with reason)

MQTT CONTROL PACKET VARIABLE HEADER (2/3)

- In PUBLISH:
 - Topic Name

MQTT CONTROL PACKET VARIABLE HEADER (3/3)

- In assorted packets:
 - Packet Identifier (2B)

Control Packet	Packet Identifier field
CONNECT	NO
CONNACK	NO
PUBLISH	YES (If QoS > 0)
PUBACK	YES
PUBREC	YES
PUBREL	YES
PUBCOMP	YES
SUBSCRIBE	YES
SUBACK	YES
UNSUBSCRIBE	YES
UNSUBACK	YES
PINGREQ	NO
PINGRESP	NO
DISCONNECT	NO

Source: *MQTT Version 3.1.1*
OASIS Standard , October 2014

MQTT CONTROL PACKET PAYLOAD

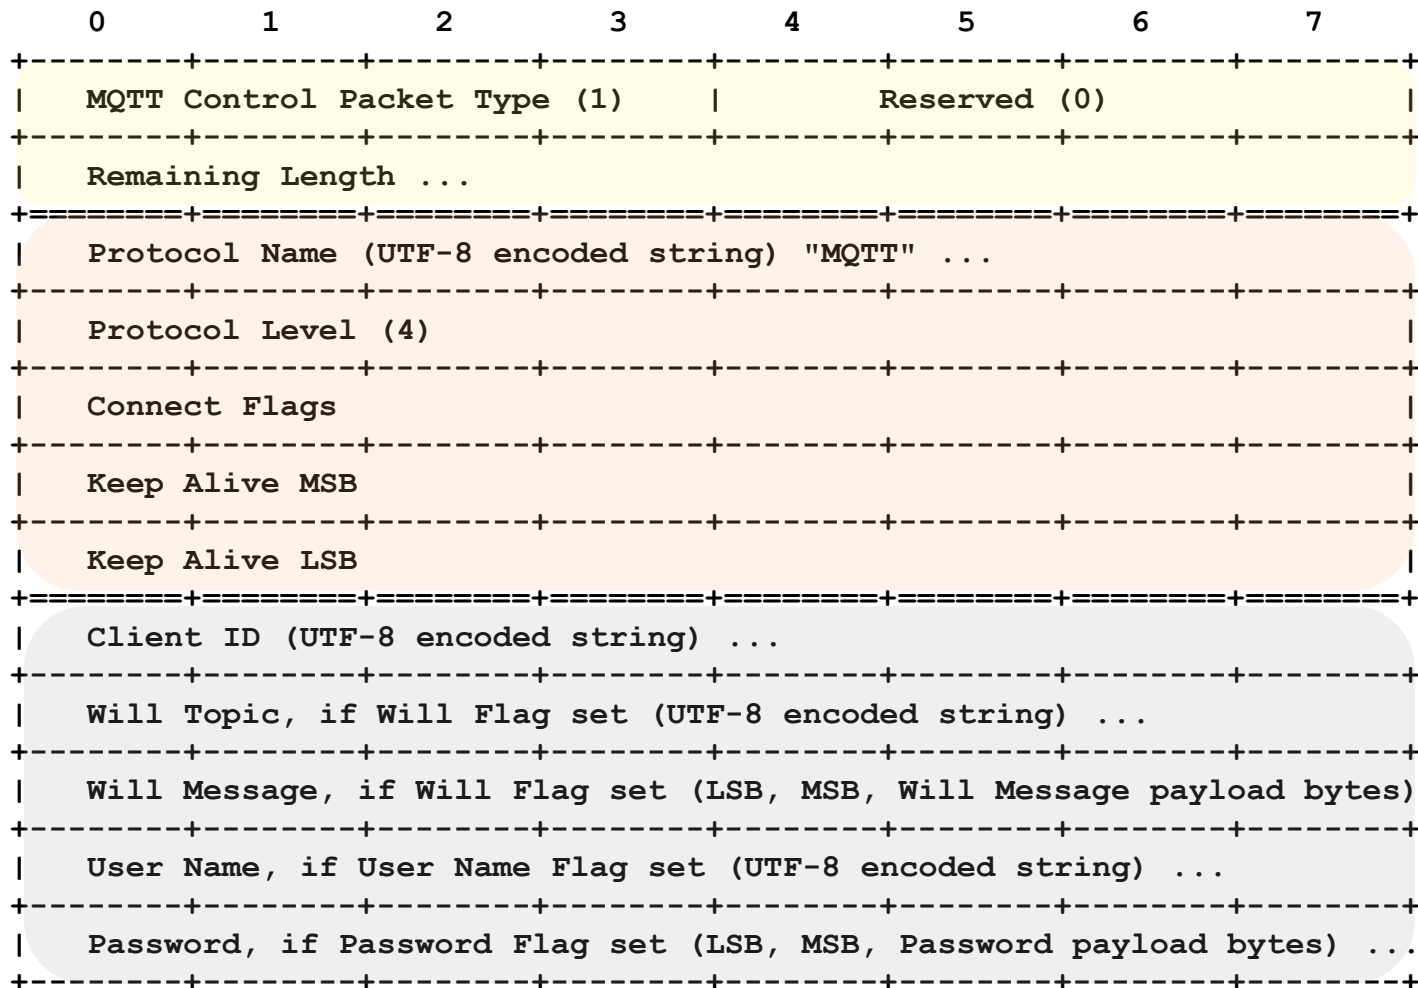
Control Packet	Payload	
CONNECT	Required	<u>ClientID</u> , Will Topic, Will Message, User Name, Password
CONNACK	None	
PUBLISH	Optional	Application Message (note: MQTT is agnostic as to the format)
PUBACK	None	
PUBREC	None	
PUBREL	None	
PUBCOMP	None	
SUBSCRIBE	Required	<u>(Topic_Filter_1,QoS_1)</u> , (Topic_Filter_2,QoS_2),...// requested QoS
SUBACK	Required	<u>QoS_1</u> , QoS_2, ... // granted QoS
UNSUBSCRIBE	Required	<u>Topic_Filter_1</u> , Topic_Filter_2, ...
UNSUBACK	None	
PINGREQ	None	
PINGRESP	None	
DISCONNECT	None	

underlined payload elements are mandatory


Figure 3.4 - Connect Flag bits

Bit	7	6	5	4	3	2	1	0
	User Name Flag	Password Flag	Will Retain	Will QoS		Will Flag	Clean Session	Reserved
byte 8	X	X	X	X	X	X	X	0

CONNECT



CONNACK


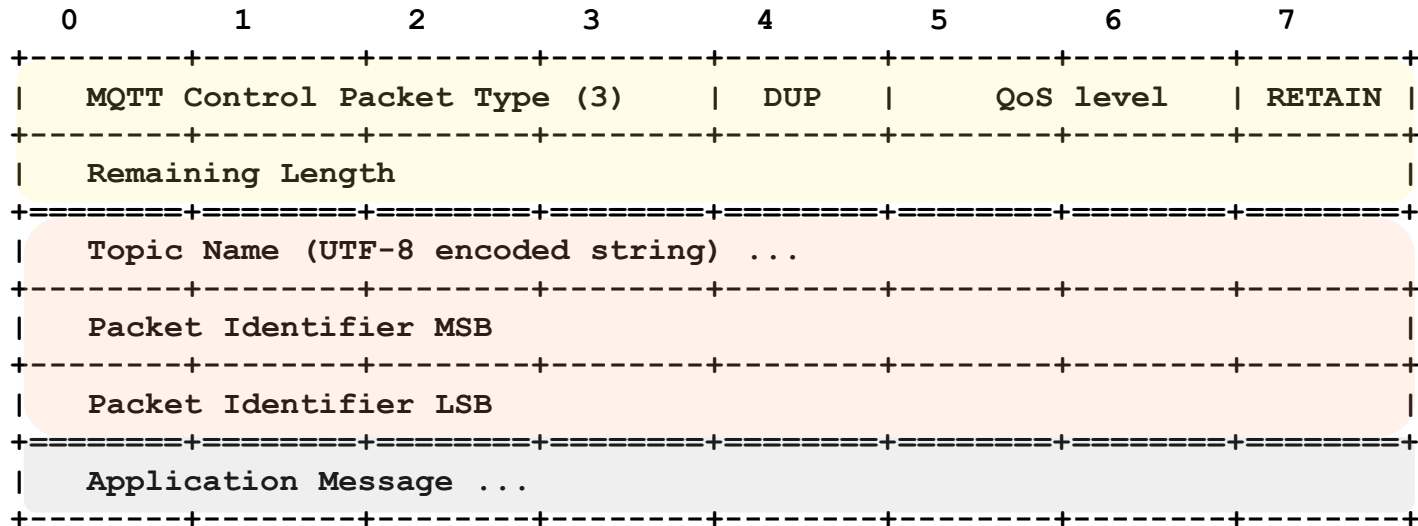


Source:
MQTT Version 3.1.1
 OASIS Standard , October 2014

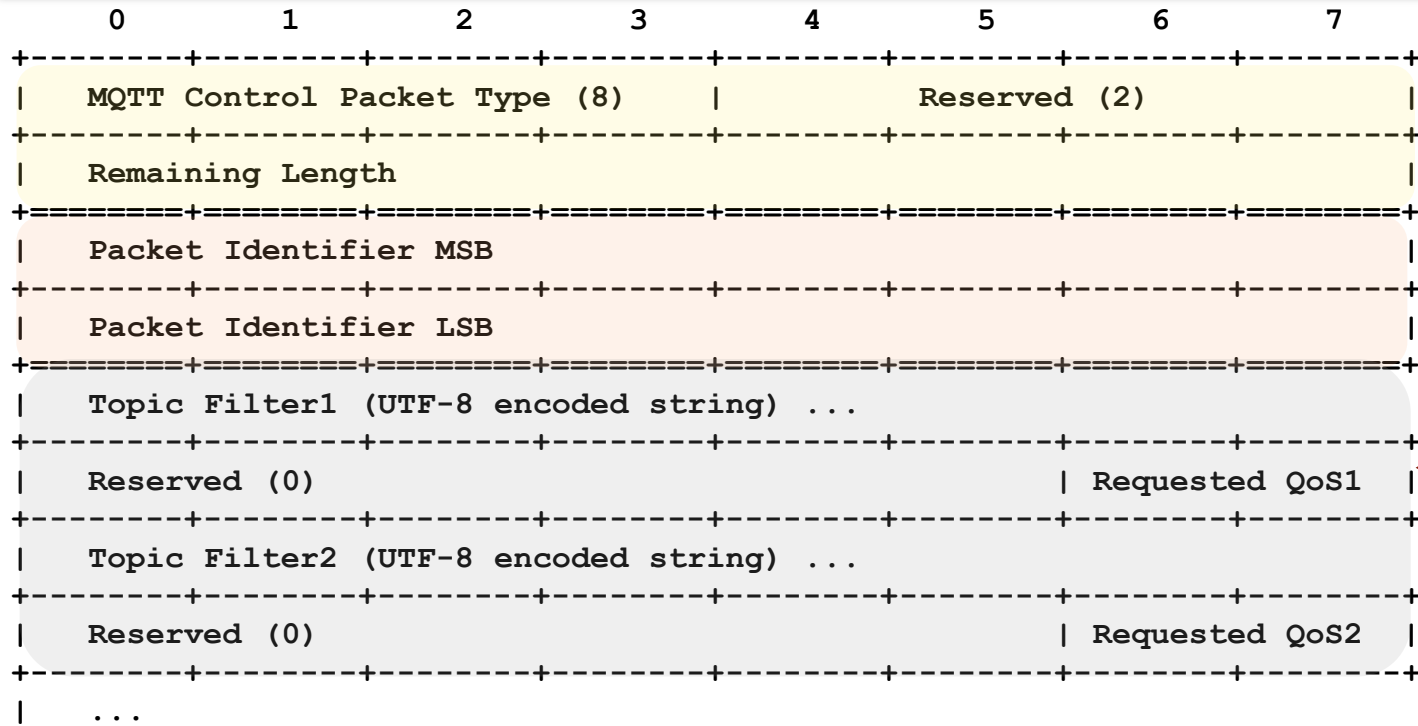
Table 3.1 – Connect Return code values

Value	Return Code Response	Description
0	0x00 Connection Accepted	Connection accepted
1	0x01 Connection Refused, unacceptable protocol version	The Server does not support the level of the MQTT protocol requested by the Client
2	0x02 Connection Refused, identifier rejected	The Client identifier is correct UTF-8 but not allowed by the Server
3	0x03 Connection Refused, Server unavailable	The Network Connection has been made but the MQTT service is unavailable
4	0x04 Connection Refused, bad user name or password	The data in the user name or password is malformed
5	0x05 Connection Refused, not authorized	The Client is not authorized to connect
6-255		Reserved for future use


PUBLISH AND PUBACK



SUBSCRIBE AND SUBACK



per topic filter:
maximum QoS
requested
("do not send at
a higher QoS")



per topic filter:
maximum QoS
granted

SUBSCRIBE AND SUBACK

- In SUBACK, the broker might grant a lower maximum QoS than the subscriber requested.
- The QoS of PUBLISH messages sent by the broker to the subscriber is equal to ...

$\min(\text{QoS of the message originally published by the publisher, maximum QoS granted by the broker}).$

TOPICS

- topics have a hierarchical structure:

topic level
↓
feit/room_121/temperature
feit/room_121/humidity
feit/room_CS300/temperature
fa/room_104/CO2
↑
topic level separator

- the broker accepts any topic, without earlier registration
 - MQTT is agnostic as to topics (just like with application messages)
- topics starting with "\$" are reserved (e.g., for monitoring purposes)

```
feit/room_121/temperature  
feit/room_121/humidity  
feit/room_CS300/temperature  
fa/room_104/CO2
```

TOPIC FILTERS

- to subscribe to more than one topic at once, use **wildcards**
 - allowed only in topic filters (in SUBSCRIBE)
- **single-level** wildcard (+)
 - example: `feiti/+/temperature` (all temperature readings from FEITI)
 - example: `+/+/co2` (all CO2 readings from entire WUT)
- **multi-level** wildcard (#)
 - example: `fa/#` (all sensor readings from FA)
 - must be the last character in the topic

TOPICS IN REAL LIFE (1/3)




Figure 1: One to One Messaging in Point-to-point Communication

Source: Amazon Web Services, *Designing MQTT Topics for AWS IoT Core*, May 2019

TOPICS IN REAL LIFE (2/3)

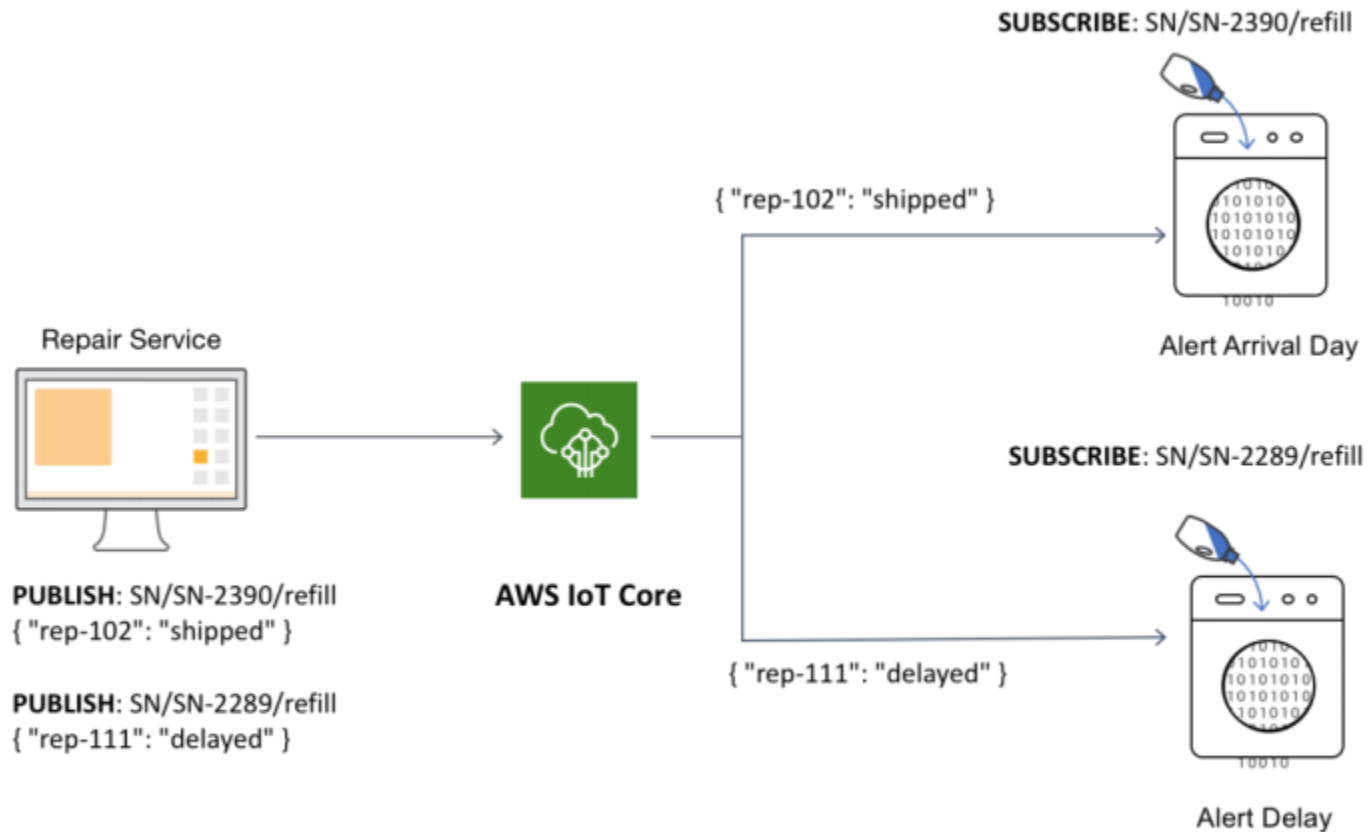


Figure 2: One-to-many messaging in point-to-point communication

Source: Amazon Web Services, *Designing MQTT Topics for AWS IoT Core*, May 2019

TOPICS IN REAL LIFE (3/3)

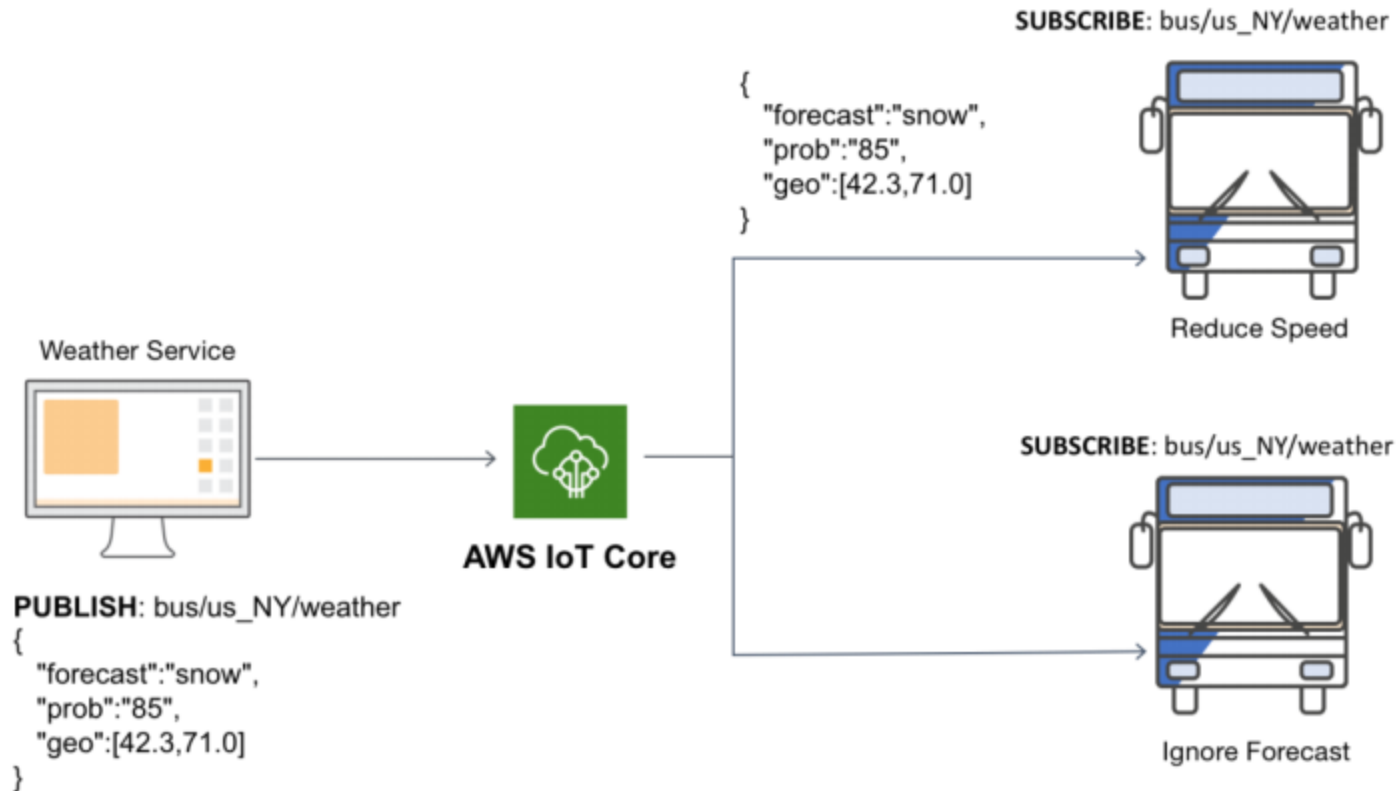



Figure 3: One-to-many messaging in broadcast communication

Source: Amazon Web Services, *Designing MQTT Topics for AWS IoT Core*, May 2019

RETAINED MESSAGES (1/2)

- a client may subscribe (to a topic) at any time
- how quickly after subscribing will the client receive the first message?
 - publishing client(s) may not have anything to publish for some time, ...
 - ... and the subscribing client does not know anything "about the topic"
- recall fixed header flags in PUBLISH



Source: *MQTT Version 3.1.1*
OASIS Standard, October 2014

- if RETAIN is set by the publisher, it is a **retained message** (i.e., stored by the broker)

RETAINED MESSAGES (2/2)

- the broker stores only one retained message per topic
 - new retained message overwrites the last one
- a client will receive a topic's retained message directly after subscribing to it
 - RETAIN will also be set, so the client knows it is a retained message
- to delete one, send a retained message with the empty payload

CONNECTIONS: KEEP ALIVE

- how will the broker find out that a client is no longer there?
 - lost connection, battery, crashing
- how will the client find out about a lost connection?
- recall the variable header part of CONNECT:
 - Protocol Name and Level, Connect Flags, **Keep Alive time** (in seconds)
- **Keep Alive time** (in CONNECT) and **keepalive timer** (at the broker)
 - for how long the server may not hear from the client before disconnection?
 - actually, it's $1.5 * \text{keep_alive_time}$
 - keepalive timer started at the broker each time a message is received
 - if no other messages to send, the client sends PINGREQ
 - the server responds with PINGRESP
 - if PINGREQ not received, the server closes the connection
 - if PINGRESP not received ("within a reasonable amount of time"), the client closes the connection

CONNECTIONS: THE SERVER'S MAY DECIDE

- "a server is permitted to disconnect a client that it determines to be inactive or non-responsive at any time, regardless of the Keep Alive value provided by that client."

Source: *MQTT Version 3.1.1*
OASIS Standard , October 2014

TESTAMENT (WILL)

Source: MQTT Version 3.1.1
OASIS Standard, October 2014

Figure 3.4 - Connect Flag bits

Bit	7	6	5	4	3	2	1	0
	User Name Flag	Password Flag	Will Retain	Will QoS		Will Flag	Clean Session	Reserved
byte 8	X	X	X	X	X	X	X	0

- recall CONNECT
 - variable header: Protocol Name and Level, Keep Alive time, **Connect Flags**
 - payload: ClientID **Will Topic, Will Message**, User Name, Password
- graceful (clean) disconnection: DISCONNECT
- when does the broker send the Will Message?
 - it has not heard from the client for $1.5 * \text{keepalive_time}$
 - the client closes the network connection without DISCONNECT
 - the broker closes the network connection because of a protocol error
 - the broker detects a network error
- the broker will send the Will Message to all that subscribed to the Will Topic

PERSISTENT SESSIONS (1/2)

- a client can lose the connection quite often
 - messages lost when not connected
 - the need to re-subscribe to topics of interest after re-connection
- **persistent session:** the broker keeps a per client session state
 - unacknowledged messages with QoS 1 or 2 (see below)
 - all subscriptions
- on a re-connection (if a client so desires) ...
 - all missed messages with QoS 1 or 2 are delivered
 - no need to re-subscribe


PERSISTENT SESSIONS (2/2)

Source: MQTT Version 3.1.1
OASIS Standard, October 2014

Figure 3.4 - Connect Flag bits

Bit	7	6	5	4	3	2	1	0
	User Name Flag	Password Flag	Will Retain	Will QoS		Will Flag	Clean Session	Reserved
byte 8	X	X	X	X	X	X	X	0

- recall CONNECT
 - variable header: Protocol Name and Level, Keep Alive time, **Connect Flags**
- Clean Session flag
 - if set, the broker discards the session state (if any) and does not maintain one
 - if not set, a persistent session is continued (or started)
 - if not set, the CONNACK SP flag indicates if the broker has a state for the client ID



← CONNACK

← Session Present Flag

QoS

- recall fixed header flags in PUBLISH
- may need reliability on top of TCP
 - a TCP connection may break down
 - there may be no space in the buffers at the receiver

- answer: application-level reliability
- QoS 0, **at most once**: QoS of the underlying network
 - no (MQTT-level) ACK
- QoS 1, **at least once**
 - the receiver responds with ACK
 - the receiver does not detect duplicates
- QoS2, **exactly once**
 - the receiver responds with ACK, etc.
 - the receiver does detect (and discards) duplicates





Table 3.2 - QoS definitions

QoS value	Bit 2	bit 1	Description
0	0	0	At most once delivery
1	0	1	At least once delivery
2	1	0	Exactly once delivery
-	1	1	Reserved – must not be used

Source: *MQTT Version 3.1.1*
OASIS Standard , October 2014

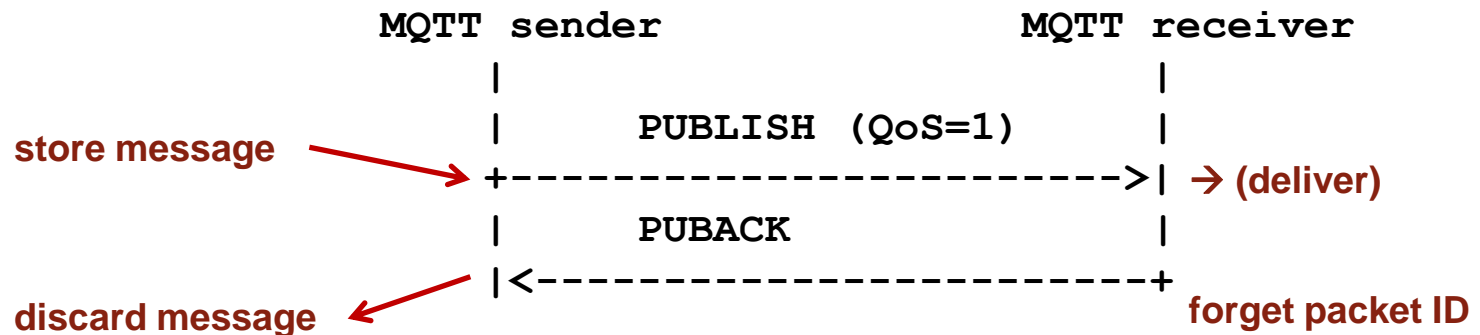
QoS 0 (AT MOST ONCE)

- fire and forget
- the message arrives once or not at all



QoS 1 (AT LEAST ONCE) (1/4)

- the receiver responds with ACK



QoS 1 (AT LEAST ONCE) (2/4)

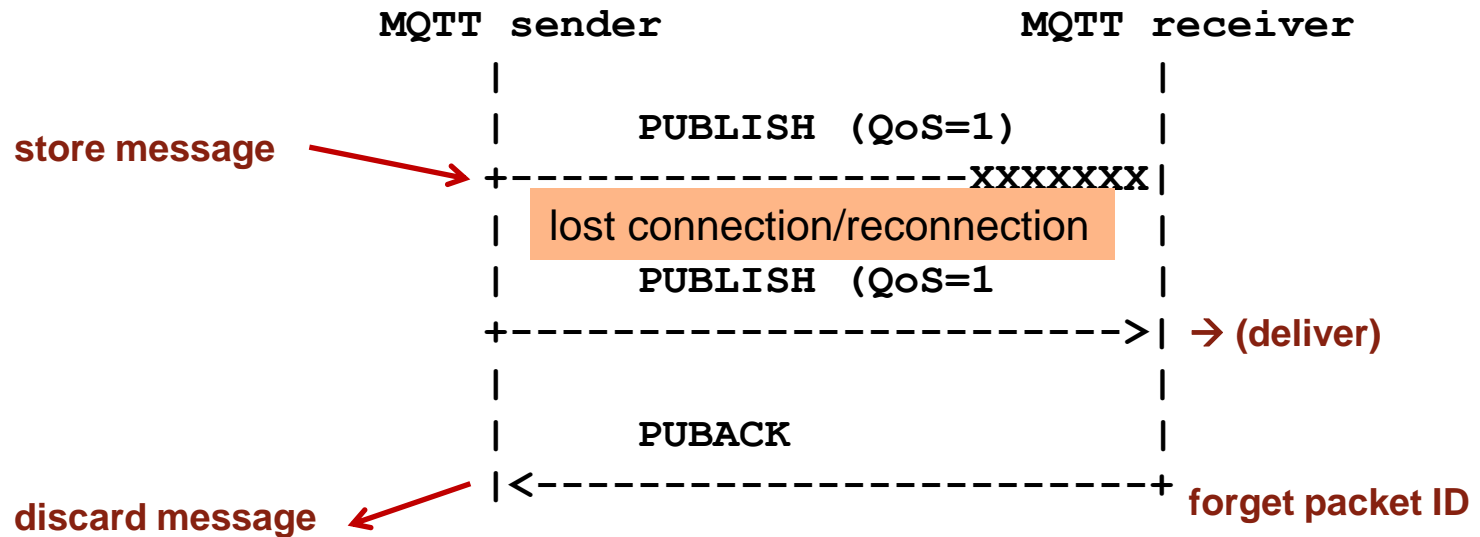
- what if either of the two messages are lost due a lost connection?
- on reconnection, the sender resends if the packet unacknowledged

”When a Client reconnects with CleanSession set to 0, both the Client and Server **MUST** re-send any unacknowledged PUBLISH Packets (where QoS > 0) using their original Packet Identifiers.

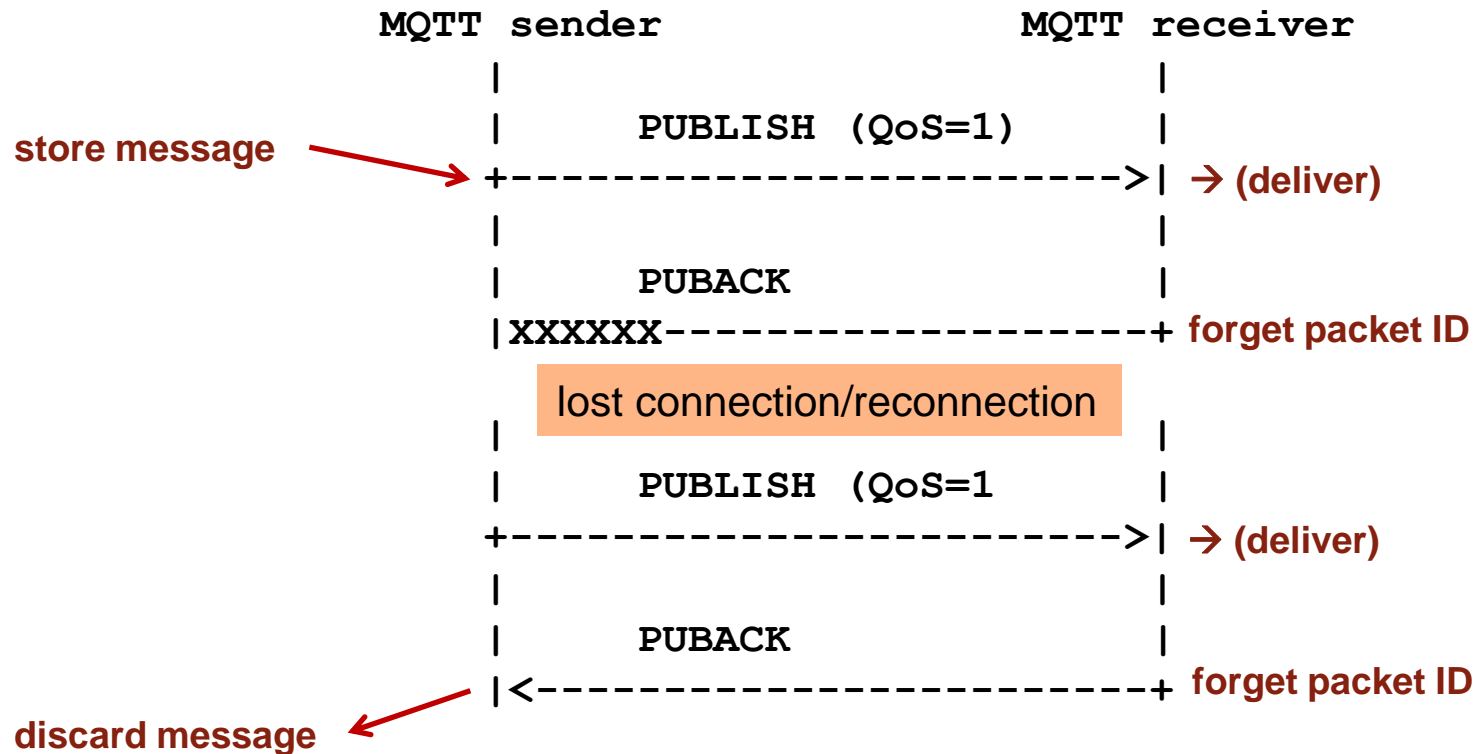
This is the only circumstance where a Client or Server is **REQUIRED** to redeliver messages.”

Source: *MQTT Version 3.1.1*
OASIS Standard , October 2014

QoS 1 (AT LEAST ONCE) (3/4)



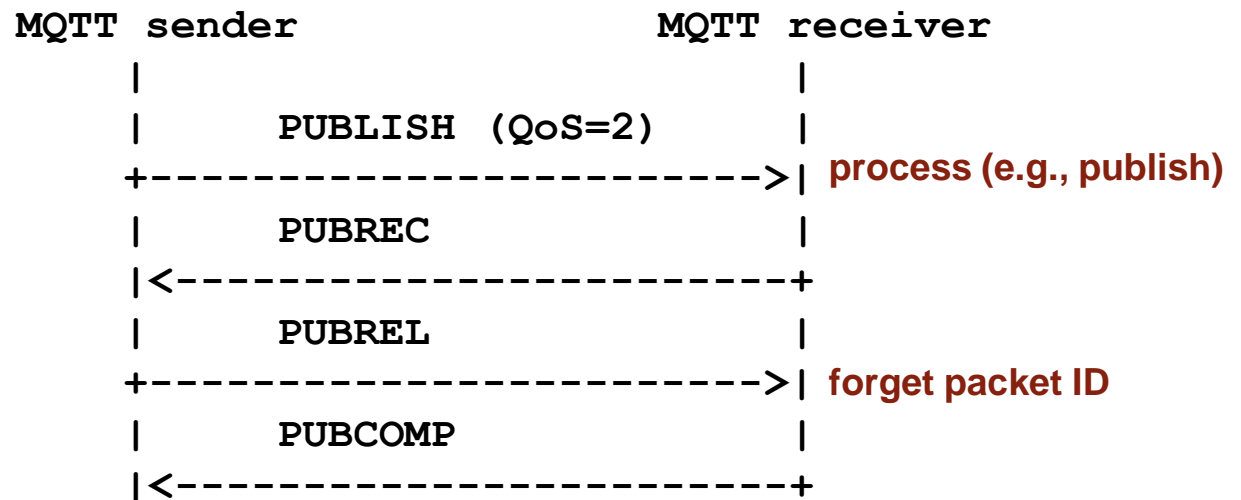
QoS 1 (AT LEAST ONCE) (4/4)



- the receiver does not detect duplicates
- multiple copies of the message may be delivered

QoS 2 (EXACTLY ONCE)

- the receiver does detect (and discards) duplicates



Thank you!

41

